

Lab 2

COSC760 - Big Data Analytics

Devere Anthony Weaver

1 Installing Scala

The following steps were used to successfully install Scala and Spark:

1. Download Scala onto Ubuntu machine

```
~ -- devere@devere-research: ~ -- ssh devere@devere-research
devere@devere-research:~$ scala
Welcome to Scala 2.11.12 (OpenJDK 64-Bit Server VM, Java 11.0.24).
Type in expressions for evaluation. Or try :help.

scala> █
```

2. Test the Scala install

```
~ -- devere@devere-research: ~ -- ssh devere@devere-research
devere@devere-research:~$ scala
Welcome to Scala 2.11.12 (OpenJDK 64-Bit Server VM, Java 11.0.24).
Type in expressions for evaluation. Or try :help.

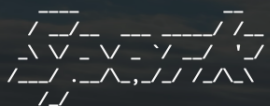
scala> println("Hello World")
Hello World

scala> █
```

3. Download Spark using `wget`
4. Start the Spark shell using `./spark-shell`

```
~ -- devere@devere-research: ~/spark-3.5.3-bin-hadoop3 -- ssh devere@devere-research
devere@devere-research:~/spark-3.5.3-bin-hadoop3$ ls
LICENSE NOTICE R README.md RELEASE bin conf data examples jars kubernetes licenses python sbin yarn
devere@devere-research:~/spark-3.5.3-bin-hadoop3$ █
```

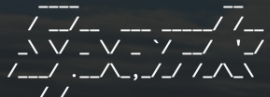
```
-- devere@devere-research: ~/spark-3.5.3-bin-hadoop3/bin -- ssh devere@devere-research ..2/screenshots
```

```
devere@devere-research:~/spark-3.5.3-bin-hadoop3/bin$ ls  
beeline find-spark-home.cmd pyspark.cmd spark-class spark-shell spark-sql.cmd spark-submit2.cmd  
beeline.cmd load-spark-env.cmd pyspark2.cmd spark-class.cmd spark-shell.cmd spark-sql2.cmd sparkR  
docker-image-tool.sh load-spark-env.sh run-example spark-class2.cmd spark-shell2.cmd spark-submit sparkR.cmd  
find-spark-home pyspark run-example.cmd spark-connect-shell spark-sql spark-submit.cmd sparkR2.cmd  
devere@devere-research:~/spark-3.5.3-bin-hadoop3/bin$ ./spark-shell  
24/10/04 15:35:55 WARN Utils: Your hostname, devere-research resolves to a loopback address: 127.0.1.1; using 192.168.1.161 instead (on i  
24/10/04 15:35:55 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address  
Setting default log level to "WARN".  
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).  
24/10/04 15:35:59 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where appli  
Spark context Web UI available at http://192.168.1.161:4040  
Spark context available as 'sc' (master = local[*], app id = local-1728070560199).  
Spark session available as 'spark'.  
Welcome to  
 version 3.5.3  
Using Scala version 2.12.18 (OpenJDK 64-Bit Server VM, Java 11.0.24)  
Type in expressions to have them evaluated.  
Type :help for more information.  
scala>
```

5. Verify Spark shell is working by running a basic command

```
-- devere@devere-research: ~/spark-3.5.3-bin-hadoop3/bin -- ssh devere@devere-research ..2/screenshots  
devere@devere-research:~/spark-3.5.3-bin-hadoop3/bin$ ls  


|                      |                     |                 |                     |                  |                  |                   |
|----------------------|---------------------|-----------------|---------------------|------------------|------------------|-------------------|
| beeline              | find-spark-home.cmd | pyspark.cmd     | spark-class         | spark-shell      | spark-sql.cmd    | spark-submit2.cmd |
| beeline.cmd          | load-spark-env.cmd  | pyspark2.cmd    | spark-class.cmd     | spark-shell.cmd  | spark-sql2.cmd   | sparkR            |
| docker-image-tool.sh | load-spark-env.sh   | run-example     | spark-class2.cmd    | spark-shell2.cmd | spark-submit     | sparkR.cmd        |
| find-spark-home      | pyspark             | run-example.cmd | spark-connect-shell | spark-sql        | spark-submit.cmd | sparkR2.cmd       |

  
devere@devere-research:~/spark-3.5.3-bin-hadoop3/bin$ ./spark-shell  
24/10/04 15:35:55 WARN Utils: Your hostname, devere-research resolves to a loopback address: 127.0.1.1; using 192.168.1.161 instead (on i...  
24/10/04 15:35:55 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address  
Setting default log level to "WARN".  
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).  
24/10/04 15:35:59 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where appli...  
Spark context Web UI available at http://192.168.1.161:4040  
Spark context available as 'sc' (master = local[*], app id = local-1728070560199).  
Spark session available as 'spark'.  
Welcome to  
 version 3.5.3  
Using Scala version 2.12.18 (OpenJDK 64-Bit Server VM, Java 11.0.24)  
Type in expressions to have them evaluated.  
Type :help for more information.  
  
scala> println("Spark shell is running")  
Spark shell is running  
  
scala>
```

2 Wordcount Program

The following commands were used to successfully implement the word count program using Spark:

1. Start the Spark shell
2. Create an RDD from a text file, count the number of words, and output the results



```

  ____
 /  _ \   \_   _/   ____
 \  __ \   /  / \   /  \
  \ ___ \  /  /  \  /   \
   \___ \_/  /___ \_/___ \
                        version 3.5.3

Using Scala version 2.12.18 (OpenJDK 64-Bit Server VM, Java 11.0.24)
Type in expressions to have them evaluated.
Type :help for more information.

[scala> val infile = sc.textFile("/home/devere/input.txt")
infile: org.apache.spark.rdd.RDD[String] = /home/devere/input.txt MapPartitionsRDD[1] at textFile
at <console>:23

[scala> val counts = infile.flatMap(line => line.split(" ")).map(word => (word, 1)).reduceByKey(_
_);
counts: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[4] at reduceByKey at <console>:23

[scala> counts.saveAsTextFile("spark_wordcount")

[scala> :q
devere@devere-research:~$
```

3. The following is the count output and the text file contents

```
~ — devere@devere-research: ~/spark_wordcount — ssh devere@devere-research
devere@devere-research:~/spark_wordcount$ cat ~/input.txt
people are not as beautiful as they
look,
as they walk or as they talk.
they are only as beautiful as they
love,
as they care as they share.
devere@devere-research:~/spark_wordcount$ cat part-00000
(talk.,1)
(are,2)
(only,1)
(as,8)
(they,7)
(love,,1)
devere@devere-research:~/spark_wordcount$ cat part-00001
(not,1)
(people,1)
(share.,1)
(or,1)
(care,1)
(beautiful,2)
(walk,1)
(look,,1)
devere@devere-research:~/spark_wordcount$ █
```

3 Import CSV Files into HIVE using PySpark

The following commands were used to successfully import a CSV file into HIVE and create a table:

1. Open Pyspark shell
2. Import a CSV file containing default data, parse the attributes, and create a tablew


```

      ____
     /  _ \   _ __   ___
    /  / \  / /  _ \  _ \
   /  /  _/ /  / ___/  / /
  /___/_\___/_/  /___/_/

version 3.5.3

```

Using Python version 3.10.12 (main, Sep 11 2024 15:47:36)

Spark context Web UI available at http://192.168.1.161:4040

Spark context available as 'sc' (master = local[*], app id = local-1728073760185).

SparkSession available as 'spark'.

```
>>> from pyspark.sql import HiveContext
```

```
>>> from pyspark.sql.types import *
```

```
>>> from pyspark.sql import Row
```

```
>>> csv_data = sc.textFile("Default.csv")
```

```
>>> type(csv_data)
```

```
<class 'pyspark.rdd.RDD'>
```

```
>>> csv_data = csv_data.map(lambda p: p.split(","))
```

```
>>> header = csv_data.first()
```

```
>>> csv_data = csv_data.filter(lambda p: p != header)
```

```
>>> df_csv = csv_data.map(lambda p: Row(Default = p[0], Student = p[1], Balance = float(p[2]), Income = float(p[3]))).toDF()
```

```
... )
```

```
>>> df_csv.show(1)
```

Traceback (most recent call last):

File "<stdin>", line 1, in <module>

AttributeError: 'PipelinedRDD' object has no attribute 'show'

```
>>> df_csv = csv_data.map(lambda p: Row(Default = p[0], Student = p[1], Balance = float(p[2]), Income = float(p[3]))).toDF()
```

```
>>> df_csv.show(5)
```

```

+-----+-----+-----+-----+
|Default|Student|      Balance|      Income|
+-----+-----+-----+-----+
|  "No" |  "No" | 729.526495207286| 44361.6250742669|
|  "No" | "Yes" | 817.180406555498| 12106.1347003149|
|  "No" |  "No" | 1073.54916401173| 31767.1389473999|
|  "No" |  "No" | 529.250604745278| 35704.4939350781|
|  "No" |  "No" | 785.655882930501| 38463.4958787229|
+-----+-----+-----+-----+

```

only showing top 5 rows

3. Check the newly created schema

```

+-----+-----+-----+-----+
|Default|Student|      Balance|      Income|
+-----+-----+-----+-----+
|  "No" |  "No" |729.526495207286|44361.6250742669|
|  "No" | "Yes" |817.180406555498|12106.1347003149|
|  "No" |  "No" |1073.54916401173|31767.1389473999|
|  "No" |  "No" |529.250604745278|35704.4939350781|
|  "No" |  "No" |785.655882930501|38463.4958787229|
+-----+-----+-----+-----+
only showing top 5 rows

>>> df_csv.printSchema()
root
 |-- Default: string (nullable = true)
 |-- Student: string (nullable = true)
 |-- Balance: double (nullable = true)
 |-- Income: double (nullable = true)

>>>

```

4. Output the new table

```

>>>
>>> hc = HiveContext(sc)
/home/devere/spark-3.5.3-bin-hadoop3/python/pyspark/sql/context.py:733: FutureWarning: HiveContext
is deprecated in Spark 2.0.0. Please use SparkSession.builder.enableHiveSupport().getOrCreate()
instead.
  warnings.warn(
/home/devere/spark-3.5.3-bin-hadoop3/python/pyspark/sql/context.py:113: FutureWarning: Deprecate
in 3.0.0. Use SparkSession.builder.getOrCreate() instead.
  warnings.warn(
>>> df_csv.write.format("orc").saveAsTable("default")

```
